

Date : 09 October 2022.

Team ID : PNT2022TMID38667;

Project Name : Early Detection of Chronic Kidney Disease Using Machine Learning.

Name : Keerthana. V (TM 2)

1.Download the dataset

The dataset was download and some changes applied in this dataset.

```
In [122...]: import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```

2.Load the dataset

```
In [123...]: k1 = pd.read_csv("abalone.csv")
```

```
In [124...]: k1.head()
```

```
Out[124]:
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

```
In [125...]: k1.columns
```

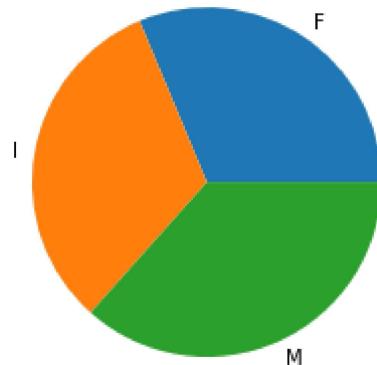
```
Out[125]: Index(['Sex', 'Length', 'Diameter', 'Height', 'Whole weight', 'Shucked weight',  
                 'Viscera weight', 'Shell weight', 'Age'],  
                 dtype='object')
```

3. perform Visualization

Univariate Analysis

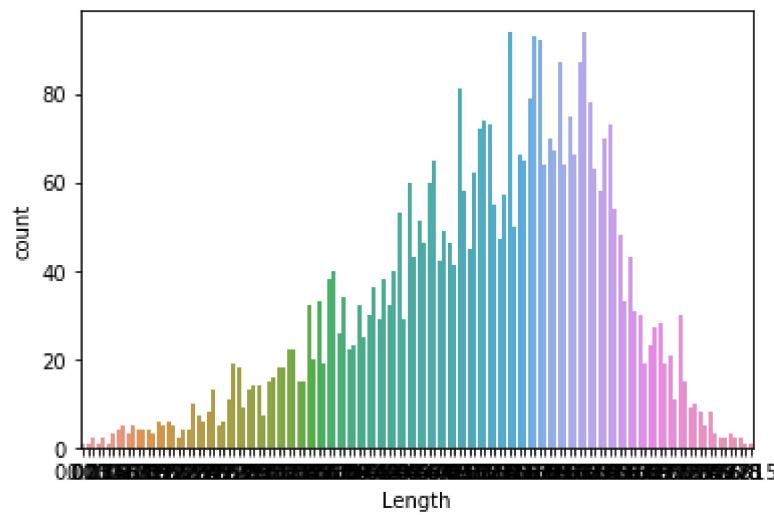
In [126...]

```
#pie chart
data1=k1.groupby("Sex",axis=0)
plt.pie(data1.count()["Length"], labels=data1.indices)
plt.show()
```



In [127...]

```
#countPlot
sns.countplot(x=k1["Length"])
plt.show()
```

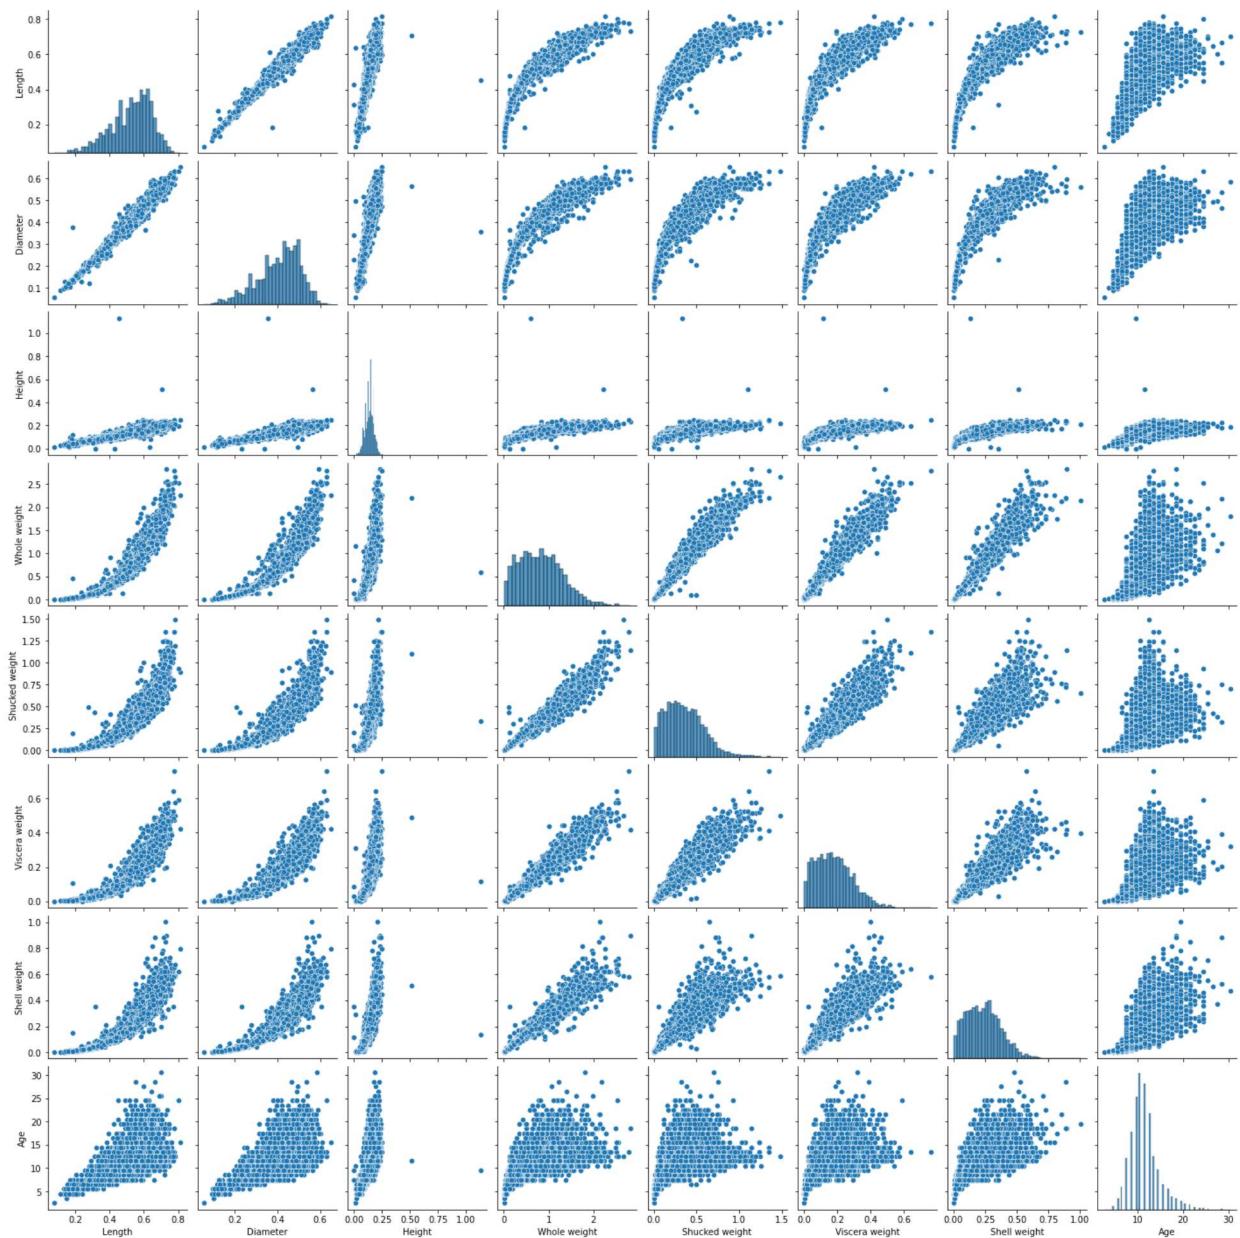


Bi-Variate Analysis

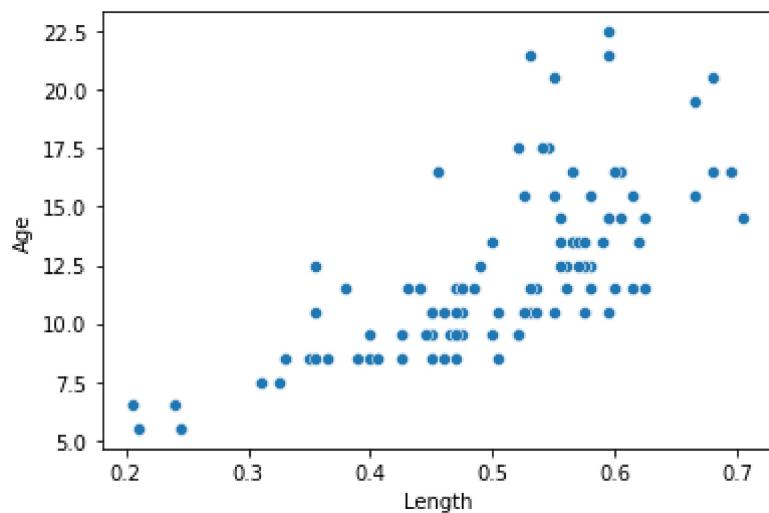
In [128...]

```
#pairPlot
sns.pairplot(k1)
plt.show()
```

Assignment 3

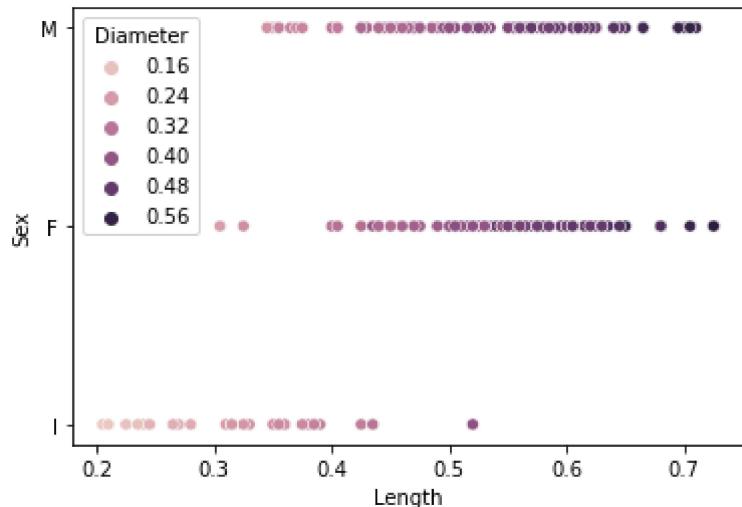


```
In [129]: #scatterplot for Length and Age
sns.scatterplot(x=k1.iloc[:100,:][["Length"]],y=f1.iloc[:100,:][["Age"]])
plt.show()
```

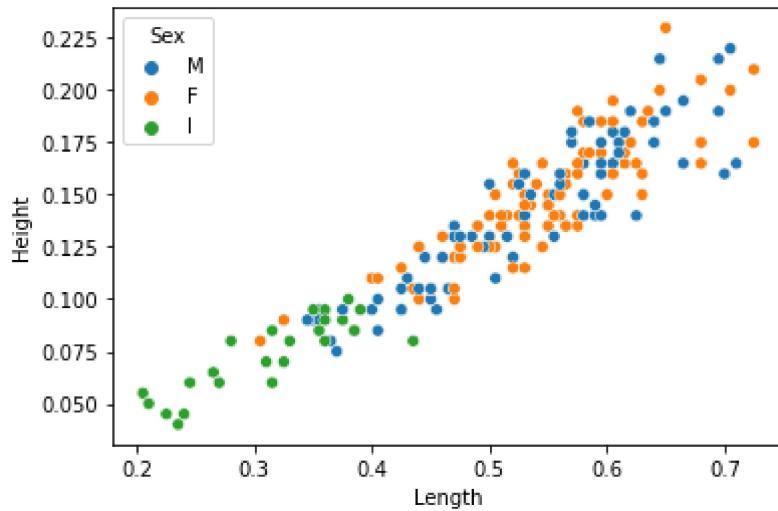


Multi-Variate Analysis

```
In [158...]: #scatterplot for Sex, Length and Diameter
sns.scatterplot(x=k1.iloc[:200,:]["Length"],y=f1.iloc[:200,:]["Sex"],hue=f1.iloc[:200,:].Diameter,plt.show()
```



```
In [159...]: #scatterplot for Length, Height and Age
sns.scatterplot(x=k1.iloc[:200,:]["Length"],y=f1.iloc[:200,:]["Height"],hue=f1.iloc[:200,:].Sex,plt.show()
```



4. Perform the descriptive statistics on the dataset

```
In [132...]: k1.describe()
```

Out[132]:

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4
mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	
std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	
25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	
50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	
max	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	

In [133...]: `k1.mode(numeric_only=True)`

Out[133]:

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Age
0	0.550	0.45	0.15	0.2225	0.175	0.1715	0.275	10.5
1	0.625	NaN	NaN	NaN	NaN	NaN	NaN	NaN

In [134...]: `k1.median(numeric_only=True)`

Out[134]:

Length	0.5450
Diameter	0.4250
Height	0.1400
Whole weight	0.7995
Shucked weight	0.3360
Viscera weight	0.1710
Shell weight	0.2340
Age	10.5000
dtype: float64	

In [135...]: `k1.skew(numeric_only=True)`

Out[135]:

Length	-0.639873
Diameter	-0.609198
Height	3.128817
Whole weight	0.530959
Shucked weight	0.719098
Viscera weight	0.591852
Shell weight	0.620927
Age	1.114102
dtype: float64	

In [136...]: `k1.kurt(numeric_only=True)`

```
Out[136]: Length      0.064621
           Diameter   -0.045476
           Height     76.025509
           Whole weight -0.023644
           Shucked weight  0.595124
           Viscera weight  0.084012
           Shell weight   0.531926
           Age          2.330687
           dtype: float64
```

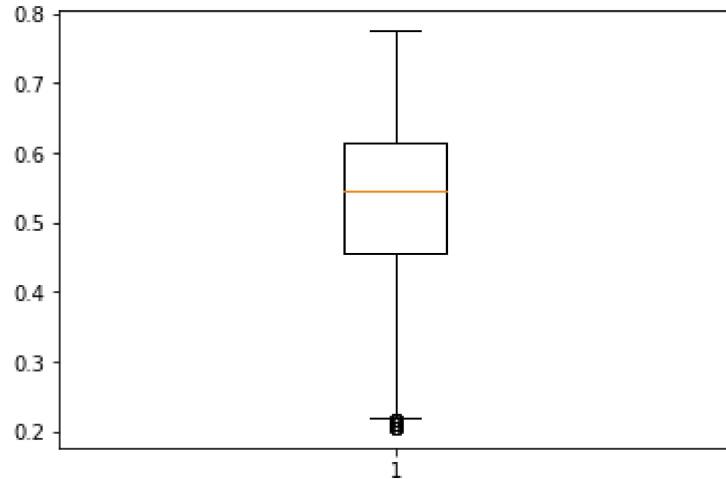
5. Handle the missing values

```
In [137... #find the null columns
k1.isnull().sum()
```

```
Out[137]: Sex      0
           Length   0
           Diameter 0
           Height   0
           Whole weight 0
           Shucked weight 0
           Viscera weight 0
           Shell weight 0
           Age      0
           dtype: int64
```

6. Find the outliers and replace the outliers

```
In [162... #find outliers
plt.boxplot(k1["Length"])
plt.show()
```



```
In [163... #handling outliers: InterQuartile Range(IQR)
Q3=np.percentile(f1["Length"],75,interpolation='midpoint')
Q1=np.percentile(f1["Length"],25,interpolation='midpoint')
IQR=Q3-Q1
print("Q1: ", Q1)
print("Q3:", Q3)
print("IQR: ", IQR)
```

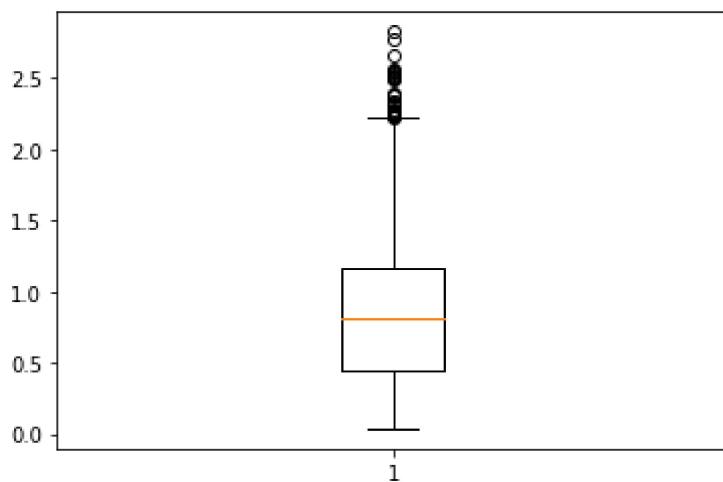
```
Q1: 0.45
Q3: 0.615
IQR: 0.1649999999999998
```

```
In [164...]: upperOutlayers=Q3+1.5*IQR
lowerOutlayers=Q1-1.5*IQR
print(upperOutlayers)
print(lowerOutlayers)
```

```
0.862499999999999
0.2025000000000004
```

```
In [141...]: k1.drop(np.where(k1["Length"]>=upperOutlayers)[0],inplace=True)
k1.drop(np.where(k1["Length"]<=lowerOutlayers)[0],inplace=True)
```

```
In [142...]: #find outliers
plt.boxplot(k1["Whole weight"])
plt.show()
```



```
In [143...]: #handling outliers: InterQuartile Range(IQR)
Q3=np.percentile(f1["Whole weight"],75,interpolation='midpoint')
Q1=np.percentile(f1["Whole weight"],25,interpolation='midpoint')
IQR=Q3-Q1
print("Q1: ", Q1)
print("Q3:", Q3)
print("IQR: ", IQR)
```

```
Q1: 0.4415
Q3: 1.153
IQR: 0.7115
```

```
In [144...]: upperOutlayers=Q3+1.5*IQR
lowerOutlayers=Q1-1.5*IQR
print(upperOutlayers)
print(lowerOutlayers)
```

```
2.22025
-0.62575
```

```
In [145...]: k1.drop(np.where(f1["Whole weight"]>=upperOutlayers)[0],inplace=True)
k1.drop(np.where(f1["Whole weight"]<=lowerOutliers)[0],inplace=True)
```

7.Check the categorical columns and perform encoding

In [146...]: `k1.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4098 entries, 0 to 4176
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
---  -- 
 0   Sex              4098 non-null    object  
 1   Length            4098 non-null    float64 
 2   Diameter          4098 non-null    float64 
 3   Height            4098 non-null    float64 
 4   Whole weight      4098 non-null    float64 
 5   Shucked weight   4098 non-null    float64 
 6   Viscera weight   4098 non-null    float64 
 7   Shell weight     4098 non-null    float64 
 8   Age               4098 non-null    float64 
dtypes: float64(8), object(1)
memory usage: 320.2+ KB
```

In [147...]: `from sklearn.preprocessing import LabelEncoder
k1["Sex"] = LabelEncoder().fit_transform(k1["Sex"])
print(k1["Sex"].unique())`

[2 0 1]

In [148...]: `k1.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4098 entries, 0 to 4176
Data columns (total 9 columns):
 #   Column           Non-Null Count  Dtype  
---  -- 
 0   Sex              4098 non-null    int32  
 1   Length            4098 non-null    float64 
 2   Diameter          4098 non-null    float64 
 3   Height            4098 non-null    float64 
 4   Whole weight      4098 non-null    float64 
 5   Shucked weight   4098 non-null    float64 
 6   Viscera weight   4098 non-null    float64 
 7   Shell weight     4098 non-null    float64 
 8   Age               4098 non-null    float64 
dtypes: float64(8), int32(1)
memory usage: 304.1 KB
```

8.Split the dataset into independent and dependent variables.

In [165...]: `X=k1.iloc[:, :-2].values #independent variables
Y=k1.iloc[:, -2].values #dependent variable`

X

```
In [165]: array([[2.      , 0.455 , 0.365 , ..., 0.514 , 0.2245, 0.101 ],
   [2.      , 0.35  , 0.265 , ..., 0.2255, 0.0995, 0.0485],
   [0.      , 0.53  , 0.42  , ..., 0.677 , 0.2565, 0.1415],
   ...,
   [2.      , 0.6   , 0.475 , ..., 1.176 , 0.5255, 0.2875],
   [0.      , 0.625 , 0.485 , ..., 1.0945, 0.531 , 0.261 ],
   [2.      , 0.71  , 0.555 , ..., 1.9485, 0.9455, 0.3765]])
```

```
In [166... Y
```

```
Out[166]: array([0.15 , 0.07 , 0.21 , ..., 0.308, 0.296, 0.495])
```

9. Scale the independent variable

```
In [151... from sklearn.preprocessing import StandardScaler
std_scaler=StandardScaler().fit_transform(X)
std_scaler
```

```
Out[151]: array([[ 1.15199292, -0.63240755, -0.4775494 , -1.12098668, -0.66806029],
   [ 1.15199292, -1.55987395, -1.54217646, -1.24510998, -1.28439797],
   [-1.27331764,  0.03006846,  0.10799548, -0.12800026, -0.31983485],
   ...,
   [ 1.15199292,  0.64837939,  0.69354036,  1.60972597,  0.74620502],
   [-1.27331764,  0.86920473,  0.80000306,  0.24436964,  0.5720923 ],
   [ 1.15199292,  1.62001086,  1.54524201,  1.36147936,  2.39653728]])
```

10. Split the data into training and testing.

```
In [152... from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size = 0.3, random_state=42)
X_train.shape
```

```
Out[152]: (2868, 5)
```

```
In [153... X_test.shape
```

```
Out[153]: (1230, 5)
```

11. Build the Model

```
In [154... #Linear Regression
from sklearn.linear_model import LinearRegression
reg= LinearRegression()
reg.fit(X_train, y_train)
```

```
Out[154]: LinearRegression()
```

12. Train the Model

```
In [155...]: y_pred = reg.predict(X_test)
```

13. Test the model And

14. Measure the performance using Metrics

```
In [156...]: from sklearn import metrics
```

```
In [157...]: print('Mean absolute error(MAE): {}'.format(metrics.mean_absolute_error(y_test, y_pred)))
print('Mean squared error(MSE): {}'.format(metrics.mean_squared_error(y_test, y_pred)))
print("Intercept: {}".format(reg.intercept_))
print('R2 Score: {}'.format(metrics.r2_score(y_test, y_pred)))
```

Mean absolute error(MAE): 0.035513348487061835

Mean squared error(MSE): 0.0027850443765074177

Intercept: -0.02261289768697705

R2 Score: 0.9420641022563094